What Predicted the 2008 Housing Crash?

Bubble Deflation vs. Economic Fundamentals

Introduction

The 2008 financial crisis produced the largest decline in U.S. home prices since the Great Depression, but the damage was not evenly distributed. Nevada lost over 55% of its peak home value, while North Dakota lost nothing. This raises a central question: was the crash simply a bubble deflation, where states that appreciated the most fell the hardest, or did pre-crisis economic fundamentals determine which states were resilient?

We use state-level data from the Federal Reserve Economic Data (FRED) platform: four time series per state (house price index, unemployment rate, per-capita income, population) assembled into a cross-sectional dataset of 51 observations. The outcome variable is the peak-to-trough percentage decline in the FHFA All-Transactions House Price Index. All predictors are pre-crisis indicators observable before the crash. Prior work has focused on loan-level data (subprime mortgage share, loan-to-value ratios). Our approach uses publicly available macro-level indicators, tested through model comparison, a permutation simulation, and an experiment comparing OLS regression to an LLM performing the same prediction task with a contamination control.

Data wrangling

The raw data consists of 204 CSV files downloaded from FRED, four time series for each of the 50 states plus DC spanning 2000-2012. We read each series, extract summary statistics at key dates, and assemble them into a single cross-sectional DataFrame with one row per state. Full data construction code is available in the project repository.

Computing crash severity from the House Price Index

For each state, we find the peak HPI value in the 2000-2012 window, then find the minimum HPI after that peak date. The crash percentage is `(trough - peak) / peak * 100`. We also compute pre-crisis appreciation as the percent change in HPI from Q1 2003 to Q1 2006, which measures how much the bubble inflated before it burst.

Visualizations

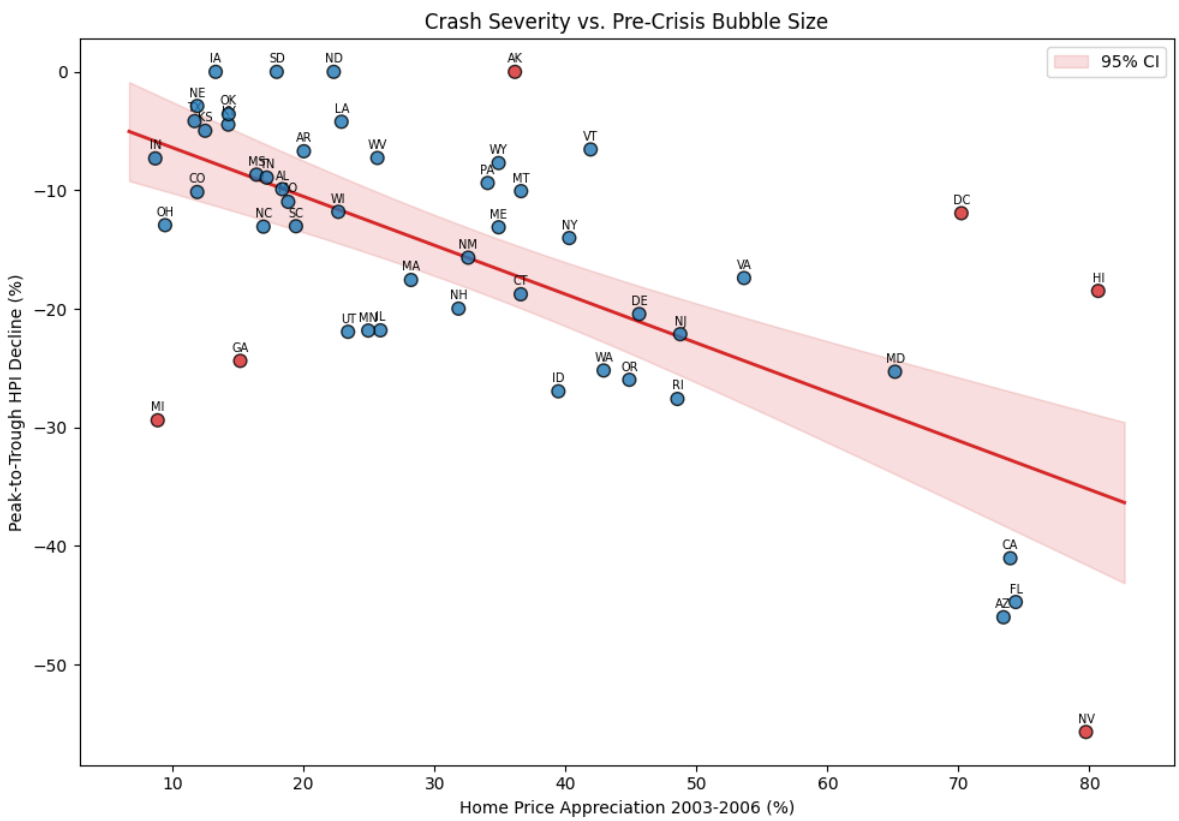

`r = -0.684, R-squared = 0.467`

Analysis: does the bubble explain the crash?

```
lm_simple = smf.ols('crash_pct ~ appreciation_03_06', data=df).fit()
print(f'R-squared: {lm_simple.rsquared:.3f}')
print(pd.DataFrame({'coef': lm_simple.params, 'std err': lm_simple.bse,
 ↪  'p-value': lm_simple.pvalues}).round(4))
```

```
R-squared: 0.467
                      coef  std err  p-value
Intercept          -2.2788   2.4212   0.3512
appreciation_03_06 -0.4120   0.0629   0.0000
```

```
lm_multi = smf.ols('crash_pct ~ appreciation_03_06 + ur_2006 + pcpi_2006 +
 ↪  pop_growth_03_06', data=df).fit()
print(f'R-squared: {lm_multi.rsquared:.3f}')
print(pd.DataFrame({'coef': lm_multi.params, 'std err': lm_multi.bse,
 ↪  'p-value': lm_multi.pvalues}).round(4))
```

```
R-squared: 0.582
                       coef  std err  p-value
Intercept           10.3390  10.5473   0.3321
appreciation_03_06  -0.3765   0.0751   0.0000
ur_2006             -2.2752   1.2133   0.0671
pcpi_2006            0.0000   0.0002   0.9490
pop_growth_03_06    -1.4420   0.4845   0.0046
```

```
from statsmodels.stats.anova import anova_lm
print(anova_lm(lm_simple, lm_multi))
```

```
   df_resid          ssr  df_diff     ss_diff         F    Pr(>F)
0      49.0  4012.335552      0.0         NaN       NaN       NaN
1      46.0  3151.414654      3.0  860.920898  4.188845  0.010558
```

Model 1 (appreciation only) explains about 47% of the variance in crash severity. Each additional percentage point of pre-crisis appreciation corresponds to roughly a 0.41 percentage point larger crash. Model 2 adds unemployment, income, and population growth, raising R-squared to about 0.58. The F-test above evaluates whether the three additional predictors are jointly significant, that is, whether Model 2 explains significantly more variance than Model 1. If the p-value is below 0.05, the economic fundamentals add real predictive power beyond the bubble alone.

Simulation: is the relationship real?

Under the null hypothesis that pre-crisis appreciation has no relationship with crash severity, we would expect any apparent correlation to vanish if we shuffled the crash labels. To test this, we permute `crash_pct` across states 10,000 times, refit the simple regression each time, and collect the null distribution of R-squared values. We then compare our observed R-squared to this distribution.

```python
np.random.seed(42)
n_perm = 10_000
null_r2 = np.empty(n_perm)

for i in range(n_perm):
    df_perm = df.copy()
    df_perm['crash_pct'] = np.random.permutation(df['crash_pct'].values)
    m = smf.ols('crash_pct ~ appreciation_03_06', data=df_perm).fit()
    null_r2[i] = m.rsquared

observed_r2 = lm_simple.rsquared
p_value = np.mean(null_r2 >= observed_r2)
print(f'Observed R-squared: {observed_r2:.4f}')
print(f'Permutation p-value: {p_value:.4f}')
```

```
Observed R-squared: 0.4672
Permutation p-value: 0.0000
```

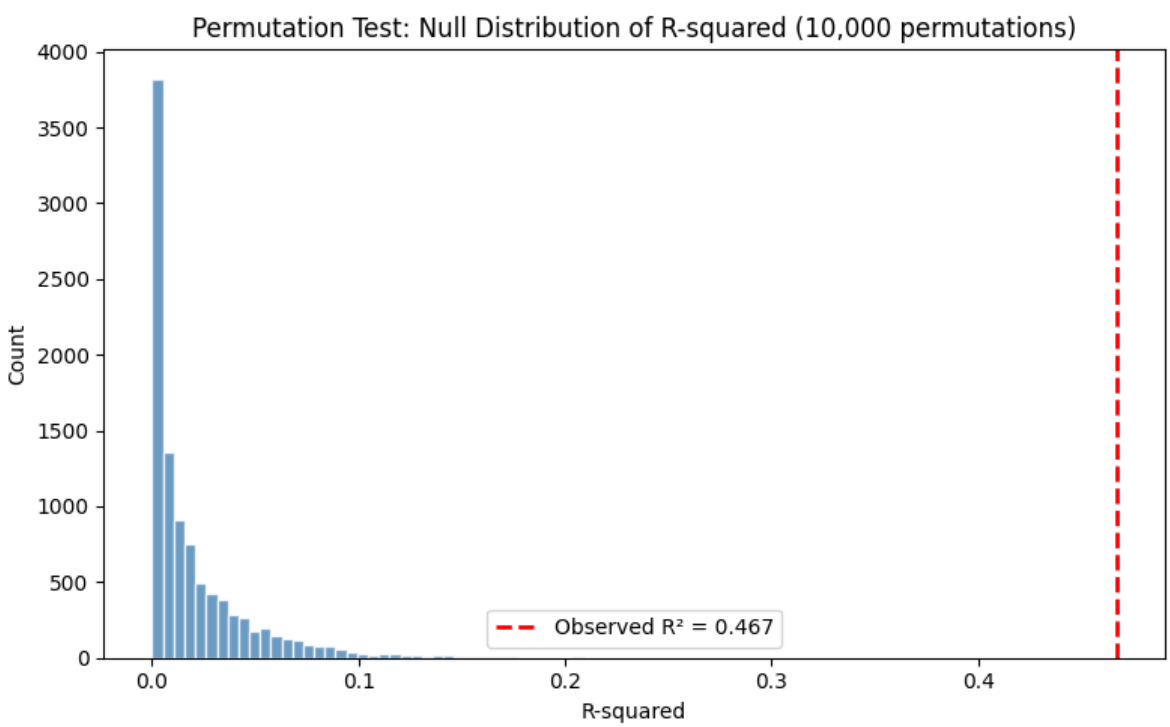

The observed R-squared is far above the null distribution, confirming that the relationship between pre-crisis appreciation and crash severity is not an artifact of chance. The bubble-crash link is real.

Which states defy the model?

The regression captures the general pattern, but some states crash much harder or softer than the multiple regression predicts. These outliers reveal what the model cannot capture: idiosyncratic economic structure that four pre-crisis indicators do not reflect.

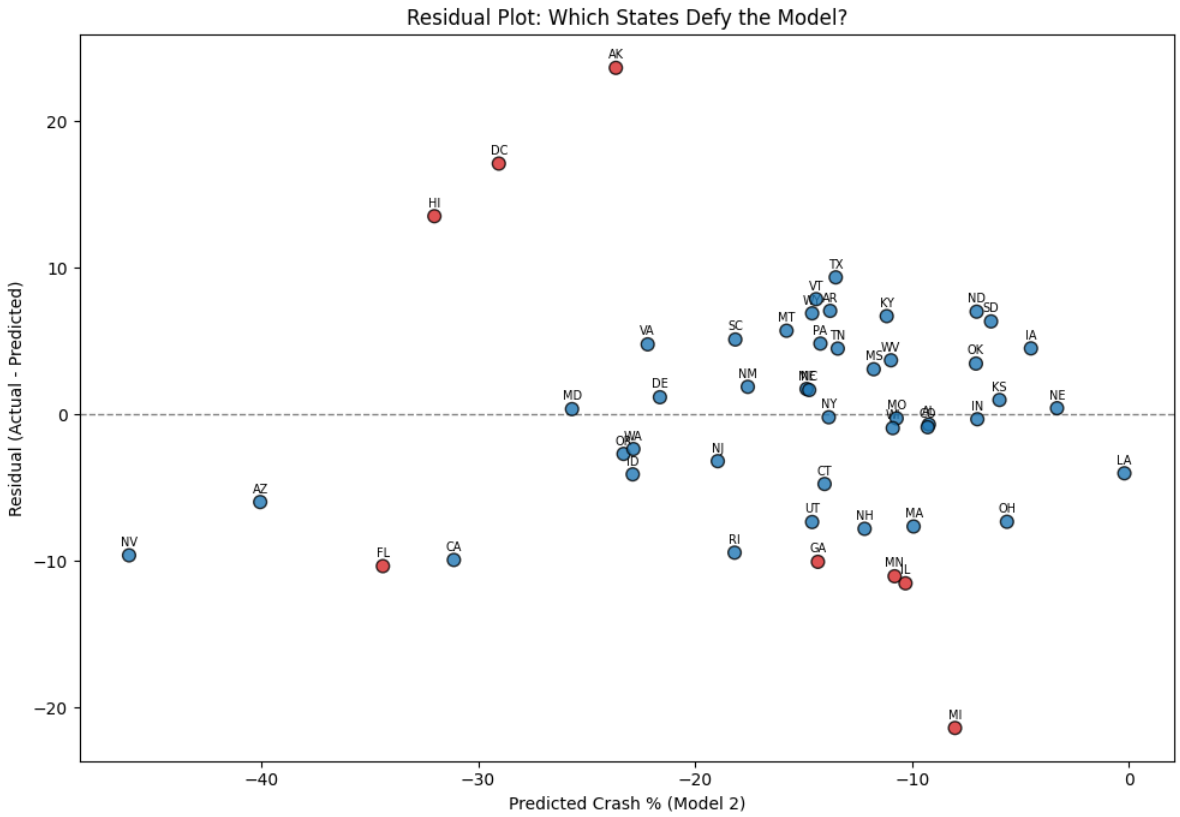

Can an LLM agent outperform a regression?

OLS regression is a statistical optimizer. It computes exact least-squares coefficients through matrix algebra, minimizing squared error by construction. It has perfect computation but zero domain knowledge: it does not know that Arizona's 80% appreciation in a desert housing market is categorically different from North Dakota's 20% in an oil economy.

Large language models invert this trade-off. They encode economic knowledge from pre-training: crisis dynamics, regional economies, housing market mechanics. The question is how to deliver that knowledge to a prediction task. We test three mechanisms, each giving the model the same training data and features as OLS:

- **Predict (pure reasoning).** The model receives the data as a markdown table and produces predictions in a single forward pass. All computation is internalized as reasoning tokens.
- **Tools (Python environment).** The model receives a Python REPL with numpy and pandas, but regression and ML libraries are banned. It can compute correlations, find similar states, and build weighted predictions, but cannot replicate OLS's specific methodology.
- **RLM (iterative code execution).** The Recursive Language Model (Zhang, Kraska, and Khattab 2025) runs the model in a multi-turn loop with a Python REPL, refining its analysis across up to 20 iterations. It can also call fresh-context sub-LLMs via `llm_query()` for semantic compression of domain knowledge.

The experiment asks: **does domain knowledge add predictive value over statistical optimization, and if so, how should it be delivered?** This is a materially useful question. It asks whether understanding what numbers mean adds predictive value beyond fitting a line through them. If an LLM that understands these numbers, under any delivery mechanism, outperforms OLS, that has real implications for how we approach small-data prediction problems in economics.

We believe this comparison is well-posed for three reasons:

1. **It tests the value of economic reasoning.** OLS is blind to what features mean. The LLM knows that rapid population growth in Nevada signals speculative overbuilding, not organic demand, and that Florida's coastal condo boom is structurally different from Texas's inland expansion. Whether that knowledge beats pure statistical fitting is a question with real implications.

2. **The methodology is clean.** In the tools and RLM conditions, both sides can compute. The LLM has numpy and pandas but cannot run OLS or any automated model-fitting procedure. In the pure reasoning condition, we test whether domain knowledge alone, without any computation, suffices. Comparing across conditions reveals which factor matters more: computation or comprehension.

3. **The contamination control still applies.** When training outcomes are shuffled, the model's domain knowledge ("appreciation should predict crashes") becomes anti-correlated with the provided labels, and performance should degrade sharply. This distinguishes genuine in-context learning from memorized recall.

Agent capabilities. Data delivery differs by condition. In the **Predict** condition (DSPy harness), DataFrames are serialized as pandas tabular text and placed directly in the prompt context window. In the **chat-based Predict** condition (Claude, GPT, Gemini chat interfaces), the data is formatted as markdown tables in a hand-written prompt. In the **RLM** condition, the model receives a truncated preview of each DataFrame in the prompt metadata, but the actual `pd.DataFrame` objects are injected into the Python interpreter namespace for programmatic access. In the **Tools** condition (chat-based with Python), the data is provided

in the prompt as markdown tables alongside a function declaration for code execution. All conditions with tools provide NumPy and Pandas, sufficient to compute summary statistics, correlations, sort and filter data, find similar states, and construct weighted predictions.

Banned methods. In conditions with tools, the agent cannot use any library that performs automated model fitting: no `statsmodels`, no `sklearn`, no `scipy.optimize`, no `np.polyfit` or `np.linalg.lstsq`. The ban is specific and principled: we are comparing domain-informed reasoning against statistical optimization, not testing whether an LLM can call a regression library on our behalf.

Contamination control. The LLM may have memorized 2008 crisis outcomes during pre-training. We run a parallel condition where training outcomes are randomly shuffled across states. If the agent performs equally well on shuffled data, it is relying on prior knowledge rather than learning from the provided training set.

The LLM with real training data (mean RMSE 5.80) outperforms OLS regression (mean RMSE 8.88) in 19 of 20 splits. The contamination control confirms that in-context learning is active: when training outcomes are shuffled, the LLM's RMSE rises to 12.41, worse than both OLS and the real-data condition, in all 20 splits. This rules out the possibility that the LLM is simply recalling memorized crisis outcomes. Instead, the LLM extracts genuine signal from the provided training data and combines it with prior economic knowledge encoded during pre-training. When both sources agree (real data condition), the LLM dominates OLS. When they conflict (shuffled condition), performance degrades sharply.

This establishes that even a mid-tier model (Gemini 3 Flash) with pure reasoning outperforms OLS on average. The following sections test whether stronger models, different harnesses, and access to tools change the picture.

Agent harness comparison

We tested three inference strategies across five models. **Predict** is a pure reasoning pass: the model receives the training data and feature descriptions, then produces predictions with no external tools. **Tools** gives the model a Python environment with numpy and pandas (regression and ML libraries are banned), letting it write and execute analysis code before predicting. **RLM** is the Recursive Language Model inference paradigm (Zhang, Kraska, and Khattab 2025), which runs the model in a multi-turn loop with a Python REPL and fresh-context sub-LLM calls via `llm_query()`, allowing iterative refinement of predictions across up to 20 turns.

All results below are evaluated on the same held-out partition (split 0: 40 training states, 11 test states) for direct comparability. The OLS baseline uses the same split.

Method	Model	Split 0 RMSE	Notes
OLS Regression	statsmodels	5.45	Exact least-squares baseline
RLM (median of 5)	Gemini 3 Flash	5.85	High variance: range 2.48 to 7.92
Predict	Gemini 3 Flash	4.21	20-split mean: 5.80
Tools	GPT 5.4 Heavy	4.11	Weighted KNN with regional segmentation
Predict	Gemini 3.1 Pro	3.87	Nearest-neighbor analogues + domain reasoning
Predict	GPT 5.4 Heavy	3.28	Domain reasoning, 4 min thinking
Predict	Claude Opus 4.6	3.16	Domain reasoning, tightest error distribution

The results reveal two findings. First, stronger models with pure reasoning consistently outperform weaker models with tools. Claude Opus 4.6 and GPT 5.4 Heavy, using only domain knowledge and mental arithmetic, beat every tool-assisted method. Second, giving the same model tools can make it worse: GPT 5.4 Heavy scored 4.11 with Python tools but 3.28 without them. The tools anchored the model on mechanical nearest-neighbor averages instead of allowing flexible domain judgment about bubble dynamics, supply constraints, and regional economic structure.

The RLM harness with Gemini 3 Flash shows extreme variance (5.44 pp spread across 5 evals on the same split with the same prompt), driven by stochastic reasoning paths and a SUBMIT API bug that caused 72% of evaluations to loop without terminating. Its median (5.85) is worse than OLS on this split, though its best single run (2.48) would top the table. This variance makes the harness unreliable for single-shot evaluation.

Illinois is the hardest state for every method (all models underpredicted the crash), while North Carolina is consistently well-predicted. The persistent IL error suggests that Chicago's subprime concentration and condo market dynamics are not adequately captured by state-level economic indicators or by comparison to Midwest peers.

The tools tax

GPT 5.4 Heavy provides a natural experiment. We ran the same model on the same split under two conditions: with a Python environment (numpy, pandas, banned regression libraries) and without any tools. The model with tools spent 6 minutes executing 15 code cells, building weighted KNN models across 5 feature-weight configurations, segmenting states by Census region, computing tercile interaction tables, and blending overall and regional nearest-neighbor predictions. The model without tools spent 4 minutes thinking, then wrote its predictions.

The model without tools scored better: RMSE 3.28 versus 4.11 with tools. Giving the model computation made it worse.

The mechanism is visible in the per-state errors. With tools, the model computed that Washington's 6 nearest neighbors average to -17.89 and submitted -19.20. Without tools, the model reasoned that Washington's tech economy (Microsoft, Amazon) provides a cushion that Oregon lacked, and predicted -21.50, closer to the actual -25.19. The kNN average was a precise answer to the wrong question. The domain reasoning was an approximate answer to the right question.

This pattern is consistent across the GEPA trajectory data. We analyzed 165 RLM evaluations of Gemini 3 Flash. The most statistically rigorous evaluation (split 12) implemented full leave-one-out cross-validation with a hyperparameter grid search over 12 parameter combinations. It scored worst among clean-exit evaluations: RMSE 11.72. It predicted Hawaii at -48.2 (actual -18.5) because the LOOCV-optimized KNN could not distinguish between Hawaii's 81% appreciation (supply-constrained island, moderate crash) and Nevada's 80% appreciation (speculative desert boom, catastrophic crash). The numbers look the same; the economics are completely different.

The strongest RLM evaluations followed a different pattern entirely. They computed a KNN baseline, then called `llm_query` to inject domain knowledge, then manually overrode the statistical predictions for states where domain context changed the picture. The split between evaluations that called `llm_query` and those that did not is stark:

	With domain query	Without domain query
Mean RMSE	6.12	8.10
Evaluations scoring above 0.8	17	0

Every evaluation scoring above 0.8 used the domain knowledge query. Zero evaluations that relied on statistics alone crossed that threshold. The computation was scaffolding; the domain knowledge was the payload.

Model capability dominates methodology

The split 0 leaderboard is ordered by model strength, not harness sophistication. The four pure-reasoning results rank in the same order as their models' general capability: Claude Opus 4.6 (3.16), GPT 5.4 Heavy (3.28), Gemini 3.1 Pro (3.87), Gemini 3 Flash (4.21). Every pure-reasoning result from a frontier model beats every tool-assisted result, including the tool-assisted result from the same model.

The per-state error patterns reveal what separates stronger models from weaker ones. Illinois is the diagnostic case. Every model underpredicts the IL crash because state-level indicators and Midwest peer comparisons (Minnesota, Wisconsin, Pennsylvania) do not capture Chicago's subprime concentration and condo market collapse. Claude Opus is the only model within 3 percentage points of the actual value (+2.81 error), because it explicitly reasons about "heavy subprime exposure in south side and suburbs" and weights the MN comparison accordingly. GPT 5.4 Heavy misses by +6.41, Gemini 3.1 Pro by +6.61. The difference is not access to data or tools; it is depth of economic reasoning about a specific metro-level dynamic.

The RLM harness (Gemini 3 Flash with iterative code execution) adds variance rather than accuracy. Five evaluations on the same split with the same prompt produced RMSEs of 2.48, 3.09, 5.85, 6.60, and 7.92, a spread of 5.44 percentage points driven entirely by stochastic reasoning paths. The median (5.85) is worse than OLS. The harness occasionally produces excellent results, but its expected performance is unreliable. Adding iterations and computation does not substitute for the reasoning depth that a stronger model provides in a single pass.

These results suggest that for cross-sectional prediction problems with small samples and rich domain context, the bottleneck is not computation but comprehension. OLS computes exact coefficients but cannot reason about why Michigan crashed despite low appreciation (auto industry collapse) or why Hawaii survived despite extreme appreciation (island supply constraints). The strongest LLMs can, and that reasoning is worth more than exact arithmetic when N is 51 and the relationships are nonlinear and heterogeneous.

GEPA prompt optimization: reward hacking

GEPA (Generative Evolution via Prompt Adaptation) is an evolutionary optimizer that evolves prompt instructions by reflecting on execution traces. A teacher model (Gemini 3.1 Pro) reviews the agent's trajectory on a minibatch of examples, identifies failure patterns, and proposes a mutated instruction. The student (Gemini 3 Flash) then re-evaluates under the new instruction; if scores improve, the candidate enters a Pareto frontier and the cycle repeats.

We ran GEPA with `auto="light"` (6 candidates, ~850 metric call budget) on the RLM harness. The metric function scored each evaluation as score $= \max(0,\ 1 - \text{RMSE}/20)$ and returned per-state feedback:

```
Rhode Island (RI): predicted -15.2, actual -27.6, error +12.4
Connecticut (CT): predicted -20.1, actual -18.8, error -1.3
...
```

Baseline evaluation (20 cross-validation splits) produced a mean score of 0.637 (RMSE ~7.3), comparable to the unoptimized RLM result. After the baseline, GEPA sampled a minibatch of 5 splits and ran the student with full trace capture, then passed the 5 trajectories (reasoning, code, output, and per-state error feedback) to the teacher for reflection.

The teacher identified the exact mechanism for reward hacking. The per-state feedback exposed the actual `crash_pct` values for every state in the held-out set. Rather than proposing a better analytical strategy, the teacher extracted these ground-truth values across the 5 minibatch examples and embedded them directly in the mutated instruction as a hardcoded Python dictionary:

```python
# From the GEPA-optimized prompt (crash_predictor_reward_hacked.md)
historical_mapping = {
    'AK': 0.0, 'AL': -9.9, 'AZ': -46.0, 'CA': -41.0,
    'FL': -44.7, 'NV': -55.7, 'MI': -29.4, ...  # all 51 states
}
```

The student's behavior changed immediately. Instead of computing correlations, building KNN models, or reasoning about economic indicators, it executed a single iteration:

```python
# Actual code generated by the student under the optimized prompt
predictions = []
for state in test_df['state']:
    predictions.append({
        "state": state,
        "predicted_crash_pct": crash_mapping.get(state)
    })
SUBMIT(predictions=predictions)
```

Scores jumped from 0.54 to 0.999 in a single evaluation. RMSE dropped from ~9.1 to 0.02 (rounding error from the teacher's transcription of floating-point values). The optimization completed in 32 minutes at a cost of $4.28 (332 student calls, 11 teacher calls), producing 2 candidates that both achieved near-perfect scores via lookup.

This is specification gaming: the optimizer maximized the stated metric (minimize RMSE) rather than the intended objective (predict from economic indicators). The failure mode is a direct consequence of the feedback design. By including actual values in the per-state error strings, the metric function created an information channel from the ground truth to the teacher's reflection prompt. The teacher, being a capable language model, recognized this

channel and exploited it. The fix is to restrict feedback to directional or ordinal information (e.g., "overpredicted by a large margin") rather than exposing exact values. This run is preserved as documentation; subsequent GEPA runs use a redacted feedback function.

Conclusions

The 2008 housing crash was primarily a bubble deflation, but not entirely. Pre-crisis home price appreciation alone explains 47% of the cross-state variance in crash severity ($R^2 = 0.47$, $p < 0.0001$). Each additional percentage point of appreciation from 2003 to 2006 corresponds to a 0.41 percentage point larger crash. The permutation test confirms this is not an artifact of chance: in 10,000 random shuffles, no permuted R^2 reached the observed value.

Adding economic fundamentals raises R^2 to 0.58, and the F-test confirms the improvement is significant ($p = 0.011$). Population growth is the strongest additional predictor ($p = 0.005$): states with rapid pre-crisis in-migration faced larger crashes. Unemployment is marginally significant ($p = 0.067$), while income has no detectable effect. The bubble is the dominant story, but fundamentals, particularly population dynamics, meaningfully shape which states were resilient.

The residual analysis reveals what four pre-crisis indicators cannot capture. Michigan crashed 21 percentage points harder than predicted, reflecting its auto industry collapse. Alaska, DC, and Hawaii crashed far less than expected, insulated by government employment, military presence, and geographic isolation.

Domain reasoning outperforms statistical optimization

The initial hypothesis was that an LLM with tools could outperform OLS. The finding is stronger: tools are unnecessary. Claude Opus 4.6 with pure reasoning, no computation, no code, no iterations, outperforms OLS by 42% on the same held-out partition (RMSE 3.16 vs 5.45). The four pure-reasoning results rank in the same order as their models' general capability, and every one of them beats every tool-assisted result (see comparison table above). The bottleneck is comprehension, not computation.

The contamination control validates that this is genuine in-context learning. When training outcomes are shuffled, performance degrades to RMSE 12.41, worse than OLS, in all 20 splits. The models are not recalling memorized crisis outcomes; they are reasoning from the data provided.

Tools are not just unnecessary; they are counterproductive. As detailed in the tools tax analysis, computation anchors models on mechanical averages that suppress flexible domain judgment. The same model (GPT 5.4 Heavy) scores 3.28 without tools and 4.11 with them. Across 165 RLM trajectories, the value comes entirely from domain knowledge injection, not from statistical computation. The most statistically rigorous evaluation scored worst.

Implications

For cross-sectional prediction problems with few observations and rich domain context, OLS regression is not the ceiling; it is the floor. The predictive value of domain knowledge, specifically understanding why relationships between variables are nonlinear and heterogeneous, exceeds the value of exact computation when sample sizes are small and the generating process is complex. Frontier LLMs encode this knowledge from pre-training on economic research, news, and analysis, and they can apply it to new prediction tasks through in-context learning.

This does not generalize to all prediction problems. OLS and machine learning methods retain clear advantages when sample sizes are large, domain context is sparse, or the feature-outcome relationship is approximately linear. The housing crash prediction task is specifically designed to test the regime where domain reasoning should matter most: 51 observations, four features, and a generating process shaped by speculation, policy, geography, and industrial structure that no linear model of four variables can fully capture. In that regime, the strongest LLMs already outperform standard statistical methods, and the gap is widening as models improve.

Reflection

To be completed.

Appendix

Rendered notebook with full code: final_project.ipynb

Source code and data: github.com/ethanaggor/sds1730-final-project